

Comprehensive Reference for Keras `tf.keras.callbacks`

This document provides an in-depth overview of the most commonly used built-in callbacks in TensorFlow/Keras. For each callback, you will find:

- **Description:** What the callback does.
- **Parameters:** List of all constructor arguments, with types, defaults, options.
- **Usage Guidance & Use Cases:** When and why to use it.
- **Pros & Cons:** Trade-offs, caveats, and best practices.

1. `EarlyStopping`

Description: Stops training when a monitored metric has stopped improving.

Constructor Parameters:

Parameter	Type	Default	Options / Notes
<code>monitor</code>	<code>str</code>	<code>'val_loss'</code>	Any model metric: <code>'loss'</code> , <code>'accuracy'</code> , <code>'val_accuracy'</code> , etc.
<code>min_delta</code>	<code>float</code>	<code>0</code>	Minimum change to qualify as improvement.
<code>patience</code>	<code>int</code>	<code>0</code>	Epochs with no improvement before stopping.
<code>verbose</code>	<code>int</code>	<code>0</code>	0 = silent, 1 = message on stop.
<code>mode</code>	<code>str</code>	<code>auto</code>	<code>min</code> , <code>max</code> , <code>auto</code> (infer from metric name).
<code>baseline</code>	<code>float</code>	<code>None</code>	Training stops if metric never exceeds this value.
<code>restore_best_weights</code>	<code>bool</code>	<code>False</code>	Restore model weights from the epoch with best metric.

Usage Guidance:

- **When to use:** Prevent overfitting and save training time by halting when no further gain.
- **Key tips:**
 - Set `monitor` to a validation metric (e.g. `'val_accuracy'`).
 - Use `restore_best_weights=True` to retain the best model.
 - Choose `patience` based on expected noise: 5–10 for stable metrics.

Pros & Cons:

- **Pros:**
 - Automatically halts training to avoid wasted epochs.
 - Can preserve best weights.
- **Cons:**
 - Too small `patience` may stop before true optimum.
 - Too large wastes compute.

2. `ModelCheckpoint`

Description: Saves model (or weights-only) at specified intervals (epochs).

Constructor Parameters:

Parameter	Type	Default	Options / Notes
<code>filepath</code>	<code>str</code>	Required	Format can include formatting options, e.g. <code>{epoch:02d}-{val_loss:.2f}.h5</code> .
<code>monitor</code>	<code>str</code>	<code>'val_loss'</code>	Metric to monitor for saving best.
<code>verbose</code>	<code>int</code>	<code>0</code>	0 = silent, 1 = message on save.
<code>save_best_only</code>	<code>bool</code>	<code>False</code>	Only save when monitored metric improves.

Parameter	Type	Default	Options / Notes
save_weights_only	bool	False	Save only model weights, not the entire model.
mode	str	auto	min , max , auto .
save_freq	str / int	'epoch'	'epoch' or integer number of samples between saves.
options	tf.saved_model.SaveOptions	None	Extra saving options.

Usage Guidance:

- **When to use:** Keep checkpoints, recover from crashes, or save best model.
- **Key tips:**
 - Use `save_best_only=True` with a validation metric to only keep the best.
 - Choose `save_weights_only=True` if you plan to rebuild architecture separately.
 - Use formatted `filepath` to track epoch and metric values.

Pros & Cons:

- **Pros:**
 - Protects against data or hardware failures.
 - Enables model selection based on validation performance.
- **Cons:**
 - Disk space usage if saving frequently or without `save_best_only` .

3. ReduceLR0nPlateau

Description: Reduces learning rate when a monitored metric has stopped improving.

Constructor Parameters:

Parameter	Type	Default	Options / Notes
monitor	str	'val_loss'	Metric to monitor.
factor	float	0.1	Factor by which to reduce LR. <code>New_lr = lr * factor</code> .
patience	int	10	Epochs with no improvement before reducing.
verbose	int	0	Print message on LR change.
mode	str	auto	min , max , auto .
min_delta	float	1e-4	Minimum change to qualify as improvement.
cooldown	int	0	Epochs to wait after LR reduction before resuming normal operation.
min_lr	float	0	Lower bound on LR.

Usage Guidance:

- **When to use:** Fine-tune learning rate schedule based on performance plateau.
- **Key tips:**
 - Set `patience` slightly larger than for `EarlyStopping` to allow adjustment.
 - Use `min_lr` to prevent LR from becoming too small.

Pros & Cons:

- **Pros:**
 - Automates LR scheduling without manual callbacks.
 - Helps model converge after plateau.
- **Cons:**
 - Complex interplay with optimizers’ built-in LR schedules.
 - Overly aggressive reduction can stall training.

4. LearningRateScheduler

Description: Schedules learning rate according to user-defined function.

Constructor Parameters:

Parameter	Type	Default	Options / Notes
schedule	function(epoch, lr) → new_lr	Required	Custom function that returns new LR.
verbose	int	0	0 = silent, 1 = print LR each epoch.

Usage Guidance:

- **When to use:** Implement custom LR decay (e.g., step decay, cosine annealing).
- **Key tips:**
 - Ensure `schedule` returns a positive float.
 - Combine with `ModelCheckpoint` or `TensorBoard` to monitor LR changes.

Pros & Cons:

- **Pros:**
 - Full control over LR over epochs.
 - **Cons:**
 - Requires careful design of schedule function.
-

5. TensorBoard

Description: Streams logs to TensorBoard for visualization.

Constructor Parameters:

Parameter	Type	Default	Options / Notes
log_dir	str	Required	Directory where to save logs.
histogram_freq	int	0	Frequency (in epochs) at which to compute activation and weight histograms.
write_graph	bool	True	Whether to log the graph.
write_images	bool	False	Write model weights to visualize as images.
update_freq	str / int	'epoch'	'batch' , 'epoch' , or integer (samples).
profile_batch	int / tuple	2	Batch(es) to profile. E.g. (10, 20) .

Usage Guidance:

- **When to use:** Visualize training metrics, graphs, histograms, embeddings.
- **Key tips:**
 - Use `histogram_freq>0` for deep introspection (slower).
 - Launch TensorBoard with `tensorboard --logdir=...`

Pros & Cons:

- **Pros:**
 - Rich visual insights into training and model internals.
 - **Cons:**
 - Writing histograms/images can slow down training.
-

6. CSVLogger

Description: Streams epoch results to a CSV file.

Constructor Parameters:

Parameter	Type	Default	Options / Notes
filename	str	Required	Path to CSV file.
separator	str	' , '	Column separator.
append	bool	False	Append to existing file if True.

Usage Guidance:

- **When to use:** Record training history for later analysis or plotting.

Pros & Cons:

- **Pros:**
 - Simple record of metrics without needing Python objects.
- **Cons:**
 - Limited to tabular logs (no histograms or images).

7. TerminateOnNaN

Description: Stops training when a NaN loss is encountered.

Constructor Parameters: None

Usage Guidance:

- **When to use:** Protect against exploding gradients or invalid loss calculations.

Pros & Cons:

- **Pros:** Immediately halts to avoid wasted compute.
- **Cons:** Doesn't provide recovery strategy—combine with `ModelCheckpoint` .

8. LambdaCallback

Description: Create custom callbacks from simple functions.

Constructor Parameters:

Parameter	Type	Default	Notes
on_epoch_begin	func	None	Called at start of each epoch.
on_epoch_end	func	None	Called at end of each epoch.
on_batch_begin	func	None	Called at start of each batch.
on_batch_end	func	None	Called at end of each batch.
on_train_begin	func	None	Called at start of training.
on_train_end	func	None	Called at end of training.

Usage Guidance:

- **When to use:** Quick ad-hoc actions (e.g. dynamic printouts, custom logging).

Pros & Cons:

- **Pros:** Lightweight, no need to subclass.
- **Cons:** Limited to simple callbacks; complex logic may merit subclassing.

Tip: For most projects, a small suite of callbacks— `EarlyStopping` , `ModelCheckpoint` , `ReduceLROnPlateau` , and `TensorBoard` — covers the majority of monitoring and control needs. Others (like `CSVLogger` or `LambdaCallback`) are useful for customized workflows

or lightweight logging.